



BITS Pilani
Pilani Campus

Advance Computer Networks (CS G525)

Virendra S Shekhawat
Department of Computer Science and Information Systems



BITS Pilani
Pilani Campus



First Semester 2018-2019

Slide_Deck_M2_2

Next...



- **SDN Controller**
 - aka Network Operating System (NOS)
- **Controller's Key Characteristics**


SDN Controller Features

- SDN Controller is a software system or a collection of systems
- It provides Management and distribution of network state
 - e.g. topology information, control session information, configuration information
- Gives *High* level data model (Yang modeling language) that captures the relationships between managed sources, policies and other services
- Gives APIs to exposes the controller services to an application

SDN Controllers



— Northbound Interface —



- Network elements has two components:
OpenFlow client,
forwarding hardware with **flow tables**.
- The SDN controller must implement the network OS functionality
 - Provide **abstraction** to the upper layer
 - Provide **control** to the underlying hardware
 - **Managing** the resources

Ten Key Characteristics of SDN Controller [1]



- OpenFlow Support
- Network Virtualization
 - VLAN (Layer-2) and VRF (Layer-3) are already there ?
 - How server virtualization is different ...?
 - Decouple virtual networks from the physical networks
- Network Functionality
 - Routing decision on multiple header fields
 - Multi-tenancy support
 - Support for adding new protocols

Ten Key Characteristics of SDN Controller [2]



- **Scalability**
 - How Layer-2 networks connect today...?
 - Through Layer-3 functionality. Multiple Layer-3 hops incur delay.
 - How SDN can solve this problem..?
 - Provides a single network view.... Hence provides better scalability
 - How many switches it can support...?
 - Depends on use cases...

Ten Key Characteristics of SDN Controller [3]



- **Performance**
 - How much time takes to setup a flow?
 - How many flows per second a controller can setup?
 - Flow setup mechanisms:
 - Proactively vs. Reactively
- **Network Programmability**
 - Different paths for inbound and out bound traffic
 - Dynamically control the traffic based on network conditions change
 - Ability to apply sophisticated filters to packets using multiple packet header fields
 - Enables programmability by implementing north bound APIs

Ten Key Characteristics of SDN Controller [4]



- **Reliability**
 - SDN controller is a single point of failure
 - Can provide quick setup of fail-over paths
 - Multipath setup
 - Design validation by controller before configuring the network
 - Supports controller Clusters

Ten Key Characteristics of SDN Controller [5]



- **Security of the Network**
 - Should support authentication and authorization of the network administrators
 - Traffic isolation for each tenant
 - Should provision for network attack detection
- **Centralized Monitoring and Visualization**
 - Flow level traffic monitoring
 - Global view of network

SDN Controller: NOX/POX

- Originally developed by Nicira
- NOX: C++ version; POX: python version
 - Nox for performance; Pox for rapid prototyping.
- POX comes with Mininet – the simulation infrastructure
- OpenFlow v.1.0
- Programming model:
 - Controller registers for events (**PacketIn**, **ConnectionUP**, etc).
 - Programmer write event handler

NOX/POX Component



NOX/POX controller

Connection
Manager

Event
dispatcher

OpenFlow
Manager

DSO
Deployer

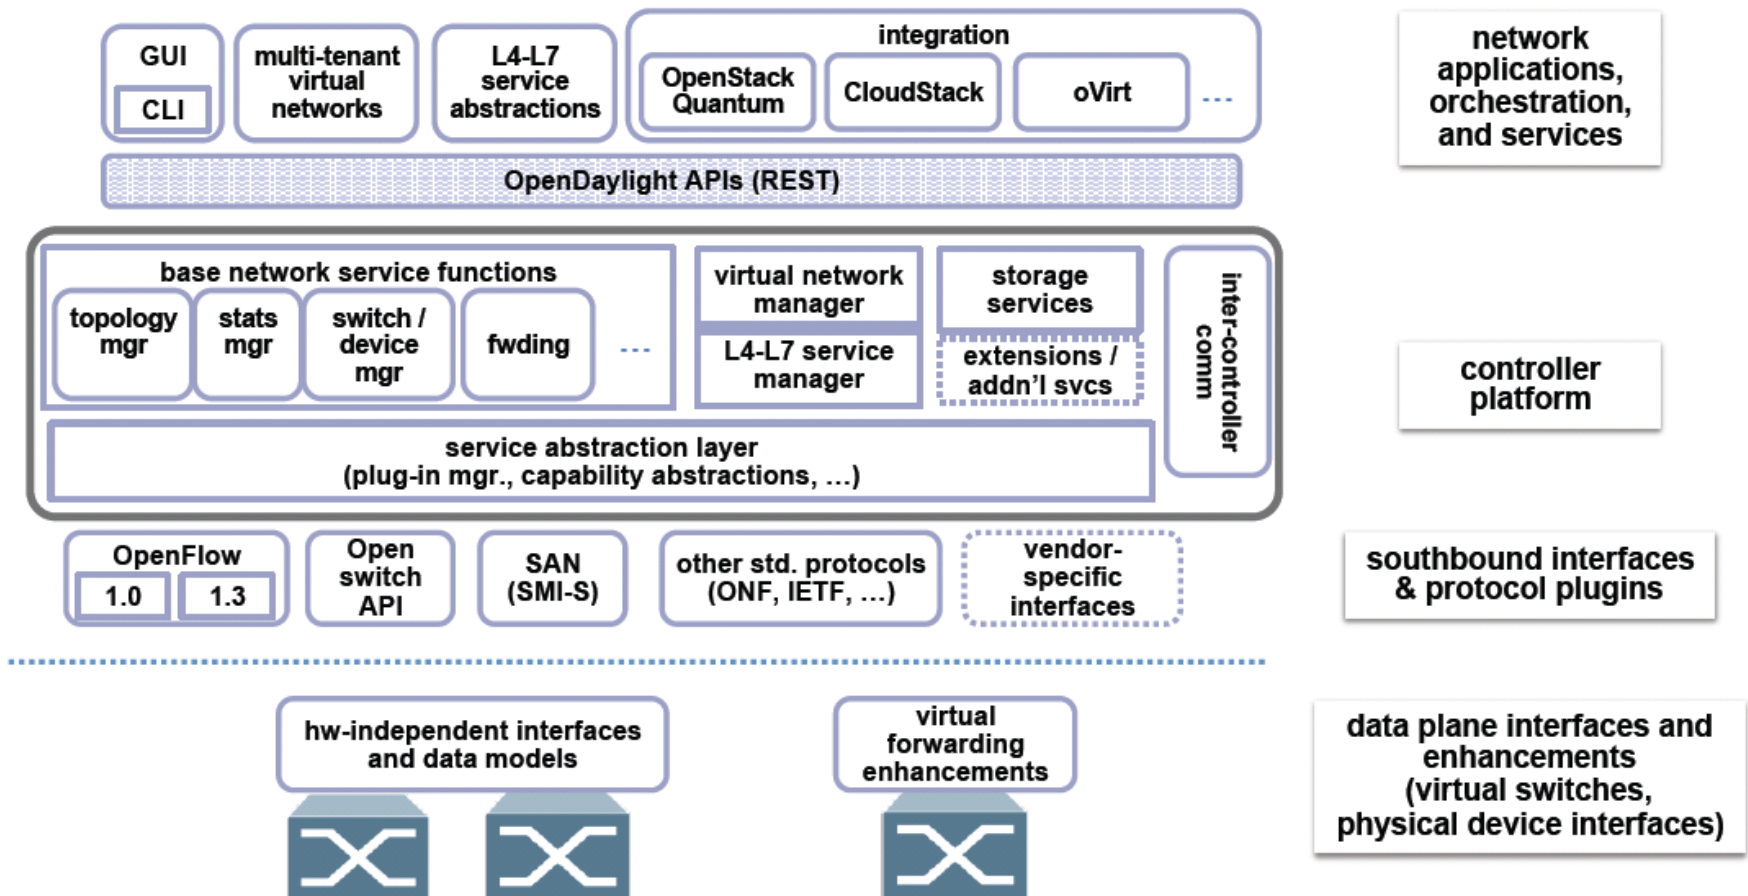
Input/output Socket
Asynchronous File

OpenFlow
API

Threading and event
management

Other
utilities

OpenDaylight Architectural Framework




Reference: <https://wiki.opendaylight.org>

Programming SDN with NOX/POX

- **Event driven programming paradigm**
 - Application writer register the events and write the event handlers.
- **API**
 - **Events:**
 - packet_in(switch, port, packet)
 - stats_in (switch, xid, pattern, packets, bytes)
 - flow_removed (switch, pattern, packets, bytes)
 - switch_join(switch)
 - Port_change(switch, port, up)
 - **Control and query openflow switches**
 - Install(switch, pattern, priority, timeout, actions)
 - uninstall(switch, pattern)
 - Query_stats(switch, pattern)

Programming SDN with POX: A Repeater Application



- Network app: when the switch is up, install rules for repeater (port 1 to port 2, port 2 to port 1)

```
def switch_join(switch):
    repeater(switch)
def repeater(switch):
    pat1 = {in_port:1}
    pat2 = {in_port:2}
    install(switch,pat1,DEFAULT,None,[output(2)])
    install(switch,pat2,DEFAULT,None,[output(1)])
```

Programming SDN

- Directly programming over NOX/POX does not have enough abstraction
 - The global network view by itself does not help
 - It is almost like assembly programming
 - Need higher level language for SDN programming
 - C++, python for Net apps.
 - Program at a higher level and then compile into the lower level API.
- Need better abstraction at the language level
 - Like all other high level languages, the abstraction needs to be sufficiently expressive for network applications
 - A good abstraction needs to have some properties that other programming language has (software reuse, etc).
 - We need to be able to implement them efficiently over NOS API.

SDN Network Updates

- Minimum updates within a single switch
- Network-wide Consistent updates

The minimum update problem

	Pattern	Priority	Action
A	[1, 2, *]	5	Port 1
B	[*, 2, 3]	4	Port 2
C	[1, *, 4]	4	Port 3
D	[1, *, 3]	3	Port 4
E	[*, *, 4]	3	Port 5
F	[*, *, 3]	2	Port 6

- Inserting rule G [*, 2, 4] in the table with high priority, above all but below rule A.
 - Insert the rule with the right priority
 - Adjust the priority of other rules if necessary
 - How to set the priority?

The minimum update problem

	Pattern	Priority	Action
A	[1, 2, *]	5	Port 1
B	[* , 2, 3]	4	Port 2
C	[1, * , 4]	4	Port 3
D	[1, * , 3]	3	Port 4
E	[* , * , 4]	3	Port 5
F	[* , * , 3]	2	Port 6

- Rule G: Pattern [* , 2, 4] above all but below Rule A.
 - Rules with overlapping patterns have interdependence
 - Changing priority would change the dependence
 - [* , 2, 4] overlaps with Rule A, C, and E.

Capturing the dependence in the flow table




- Finding the exact dependence can be tricky:
 - Two rules may be independent even if their patterns overlap.
 - Two rules are directly independent iff the patterns are disjoint or if the joint of the pattern is shadowed by other rules they both depend on.

The minimum update problem



- Represent the minimum dependency in a flow table with a DAG

	Pattern	Priority	Action
A	[1, 2, *]	5	Port 1
B	[*, 2, 3]	4	Port 2
C	[1, *, 4]	4	Port 3
D	[1, *, 3]	3	Port 4
E	[*, *, 4]	3	Port 5
F	[*, *, 3]	2	Port 6



- Insert G [*, 2, 4] into the DAG

Next...

- Retrospective on Evolving SDN
- Reading
 - Fabric: A Retrospective on Evolving SDN [Martin Casado, 2012]

Background [1]

- **Network Infrastructure**
 - Hardware
 - Software (Controls the overall behavior of the network)
- **Hardware for Ideal Network Design would be:**
 - Simple (inexpensive to build and operate)
 - Vendor-neutral
 - Future proof
- **Ideal software (control plane) for Network**
 - Flexible
 - Supports isolation, virtualization, Traffic engineering, access control and support future requirements

Background [2]

- Question

- Is today's network (aka conventional network) infrastructure *ideal*....???
- What are the inadequacies...?

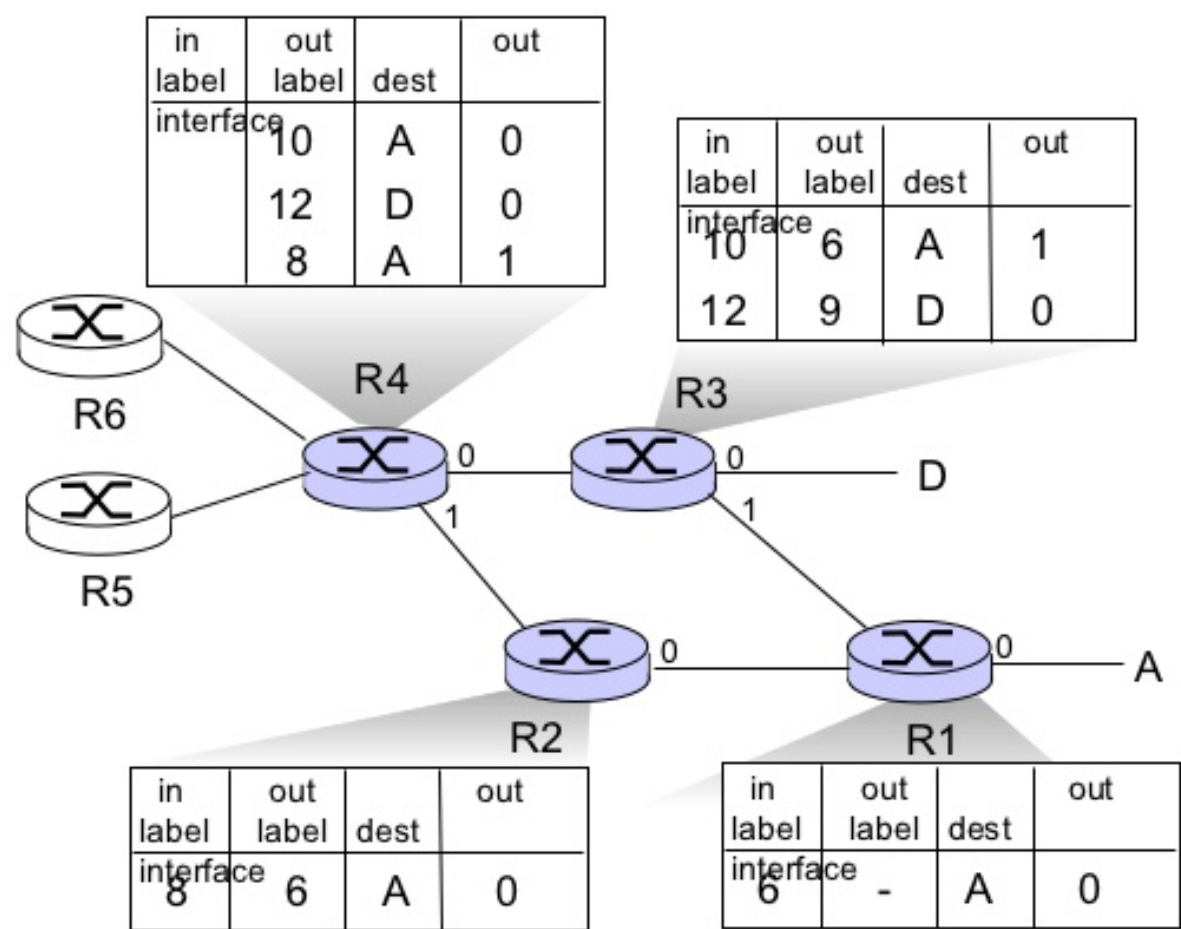
- Observation

- Inadequacies in these infrastructural aspects are more problematic than Internet's architectural deficiencies

Attempts Made So far...

- Active Networking
 - Flexibility vs. Practicality
- ATM
 - Complicated approach... Why?
- MPLS
 - VPN deployment and Traffic Engineering
- Question
 - Does MPLS meet all goals of an ideal network?

MPLS Forwarding Example



View of Network

- Network requirements comes from two sources
 - Hosts (aka users) and Operators
- Interface based view of Network
 - **Host – Network**
 - How hosts inform the network about their requirements?
 - e.g., L3 header tells about the destination address
 - **Operator – Network**
 - How operators inform the network about their requirements?
 - Usually on per group of packets (manual configuration). How SDN does???
 - **Packet – Switch**
 - How a packet identifies itself to a switch?
 - e.g. A router use some fields from the packet header as an index to its forwarding table
- Question
 - How these interfaces visualize/realize in the Internet?

Interface view of MPLS

- **How MPLS looks like?**
 - It does explicit distinction between the **network edge** and the **network core**.
 - Create traffic tunnels between edge router pairs to meet TE requirements and uses labels to forward the traffic in the network core
- **Question**
 - How these interfaces visualize/realize in MPLS network?
- **Answer**
 - Host specifies its requirement to the network through IP and packet express its requirement to the switch by MPLS label
 - MPLS distinguishes between *Host – Network* and *Packet – Switch* interfaces
 - *Operator – Network* Interface is missing?


Interface view of SDN

- Provides fully programmatic *Operator – Network* interface
 - Achieves this flexibility by decoupling control plane from the topology of the data plane
- OpenFlow Switch
 - Exports an interface that allows a remote controller to manage its forwarding state via flow tables
- Question
 - What about *Host– Network* and *Packet – Switch* interface?

Problems with SDN

- OpenFlow switch hardware is not simple
 - Requires to support lookups over so many bits
- Does not provide sufficient flexibility
 - Host requirements to continue to evolve, this leads to increase in generality for *Host-Network* interface
 - This generality must be present on every switch
- It couples the host requirements to the network core behavior
 - Change in external network protocols necessitates a change in the matching behavior (e.g. IPv4 to IPv6)

Extending SDN [1]



- **Fabric**
 - Collection of forwarding elements, provides basic packet transport
- **Edge**
 - Responsible for complex network services

Edge and Fabric Separation

- **Separation of Forwarding**
 - Allows independent evolution of fabric and edge
 - Fabric should provide a minimal set of forwarding primitives without exposing any internal details
 - External address should not be used in forwarding decisions
- **Separation of Control**
 - Two different problems
 - Fabric is responsible for packet transport across the network
 - Edge is supposed to provide more semantically rich services like security, mobility, isolation etc.
 - Separation allows any Fabric can support any edge!!!

Addressing and Forwarding

- **Fabric Path Setup**
 - Use standard **IGP and ECMP**
 - Works well in **Intra-domain** network
 - Use **MPLS**
 - Works well in **Inter-domain** network
- **Addressing and Forwarding in the Fabric**
 - Fabric forwarding elements are different from traditional forwarding elements... How?
 - Two approaches can be used-
 - Fabric addresses are opaque labels and does MPLS like forwarding (**Suitable for both path based or destination based with label-aggregation**)
 - Can follow destination based lookup with longest prefix match

Mapping the edge context to the fabric



- When a packet crosses from the edge to the fabric, something in the network must decide with which fabric-internal network address to associate with the packet
 - Two possible ways to do this-
 - Address translation
 - Edge addresses are replaced with fabric internal addresses and translated back into appropriate edge addresses at the destination
 - *Drawback:* Unnecessary coupling between both addresses!
 - Encapsulation
 - Packet is encapsulated with the another header that carries the fabric-internal addresses
 - At destination, this header is removed

- Isn't this just another approach to layering?
 - *Edge and Core are now two different Layers*
 - *Horizontal layering*
 - *Host-Network interface occurs at the edge*
 - *Packet-Switch interface occurs in the core*
- How OpenFlow will be affected?
 - “Edge” and “Core” version of the OpenFlow
 - Edge version can be implemented in the “software” like NFV

Thank You!