



BITS Pilani
Pilani Campus

Advance Computer Networks (CS G525)

Virendra S Shekhawat
Department of Computer Science and Information Systems



BITS Pilani
Pilani Campus

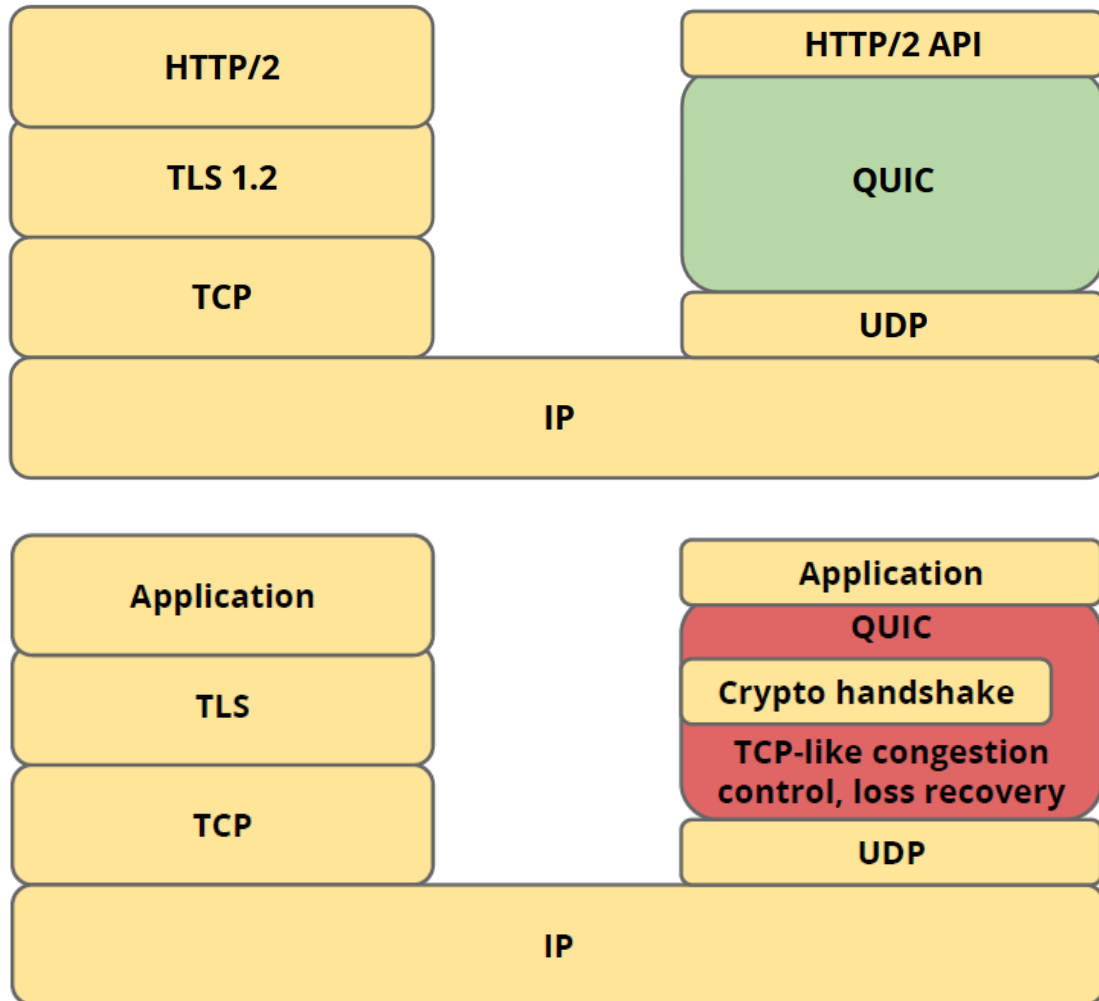


First Semester 2018-2019

Slide_Deck_M3_2

- **Quick UDP Internet Connections**
 - A new transport protocol for the internet, developed by Google.
 - Similar to TCP+TLS+HTTP2, but implemented on top of UDP
 - Faster connection establishment than TLS/TCP
 - Deals better with packet loss than TCP
 - Has Stream-level and Connection-level Flow Control
- **Reference**
 - The QUIC Transport Protocol [Adam 2017]

Where does it fit?



HTTP/2 vs HTTP/1



HTTP/2	HTTP/1
Pipelining - Parallel requests made asynchronously but server responses synchronously	Multiplexing - Multiple asynchronous HTTP requests over single TCP connection
NA	Server Push – Multiple responses for a single HTTP request
Handles multiple requests (multiplexed streams) over a single TCP connection	One outstanding request per TCP connection
Allows the server to send additional cacheable information (responses) to the client	One response per HTTP request

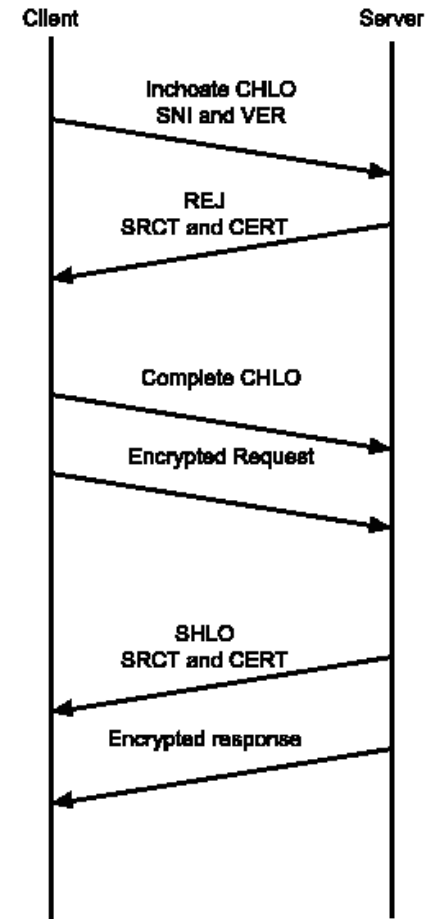
*Bi-directional sequence of text format frames sent over the HTTP/2 protocol exchanged between the server and client are known as “streams”.

QUIC Design Rationales


- Deployability and evolvability
- Low latency connection establishment
- Multi-streaming
- Better loss recovery and flexible congestion control
- Resilience to NAT-rebinding (Connection IDs vs. 4-tuple)
- Multipath for resilience and load sharing

First Ever Connection -1 RTT

- **First CHLO is inchoate (empty)**
 - Simply includes version and server name
- **Server responds with REJ**
 - Includes server config, certificates, etc.
 - Allows client to make forward progress
- **Second CHLO is complete**
 - Followed by initially encrypted request data
- **Server responds with SHLO**
 - Followed immediately by forward-secure encrypted response data



Connection Establishment

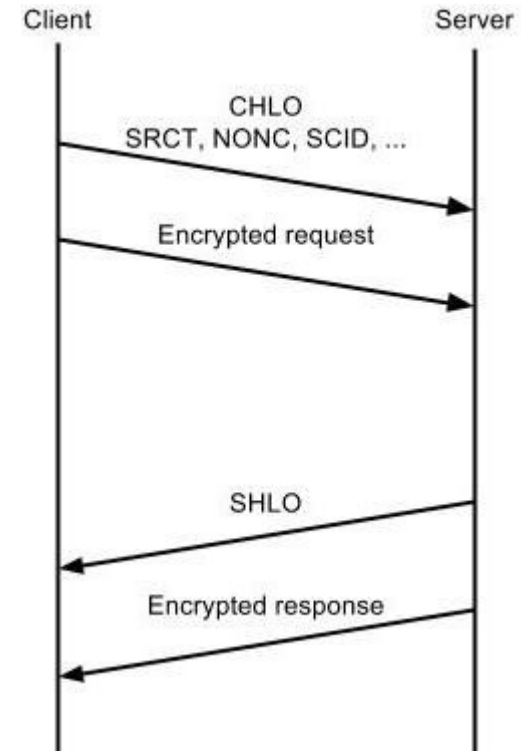


- TCP handshake
- Data
- TLS v1.2 setup

Source modified from: Chromium blog, 2015

Subsequent Connections

- **First CHLO is complete**
 - Based on information from previous connection
 - Followed by initially encrypted data.
- **Server responds with SHLO**
 - Followed immediately by forward-secure encrypted data


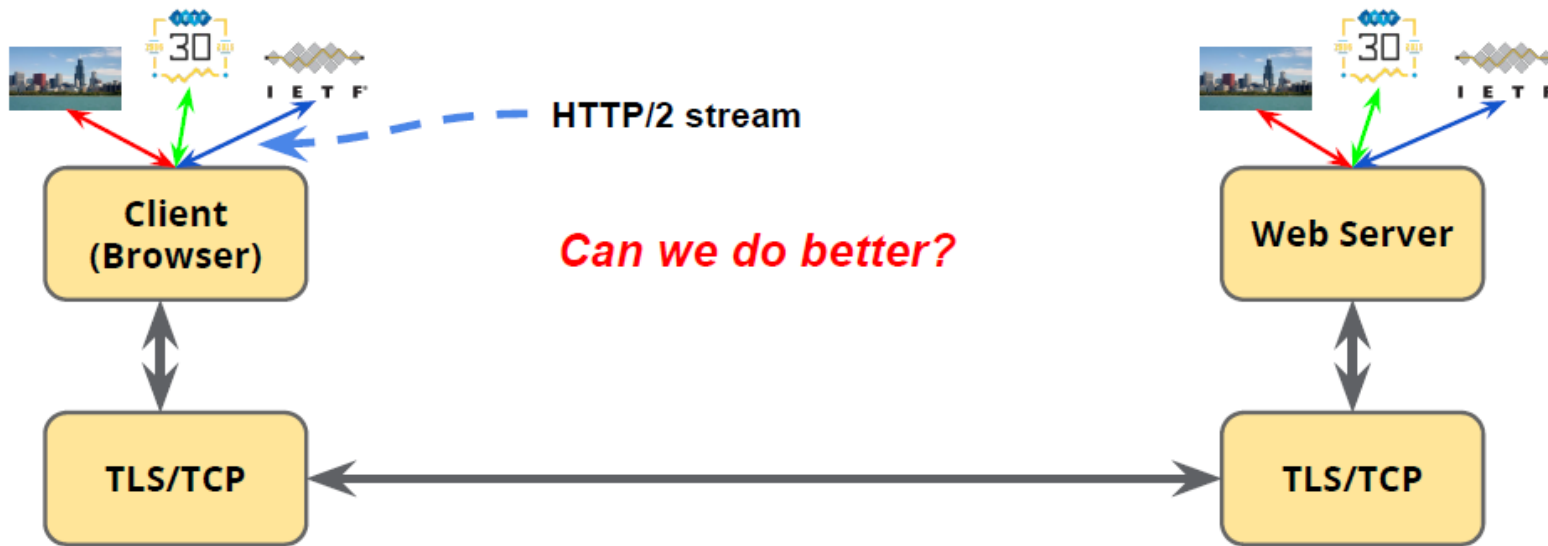


Congestion Control and Reliability



- **QUIC builds on decades of experience with TCP**
 - QUIC has pluggable congestion control
- **Retransmitted packets consume new sequence number**
 - No retransmission ambiguity
 - Prevents loss of retransmission from causing RTO
- **More verbose ACK**
 - TCP supports up to 3 SACK ranges
 - QUIC supports up to 256 NACK ranges
 - Per-packet receive times, even with delayed ACKs
- **ACK packets consume a sequence number**

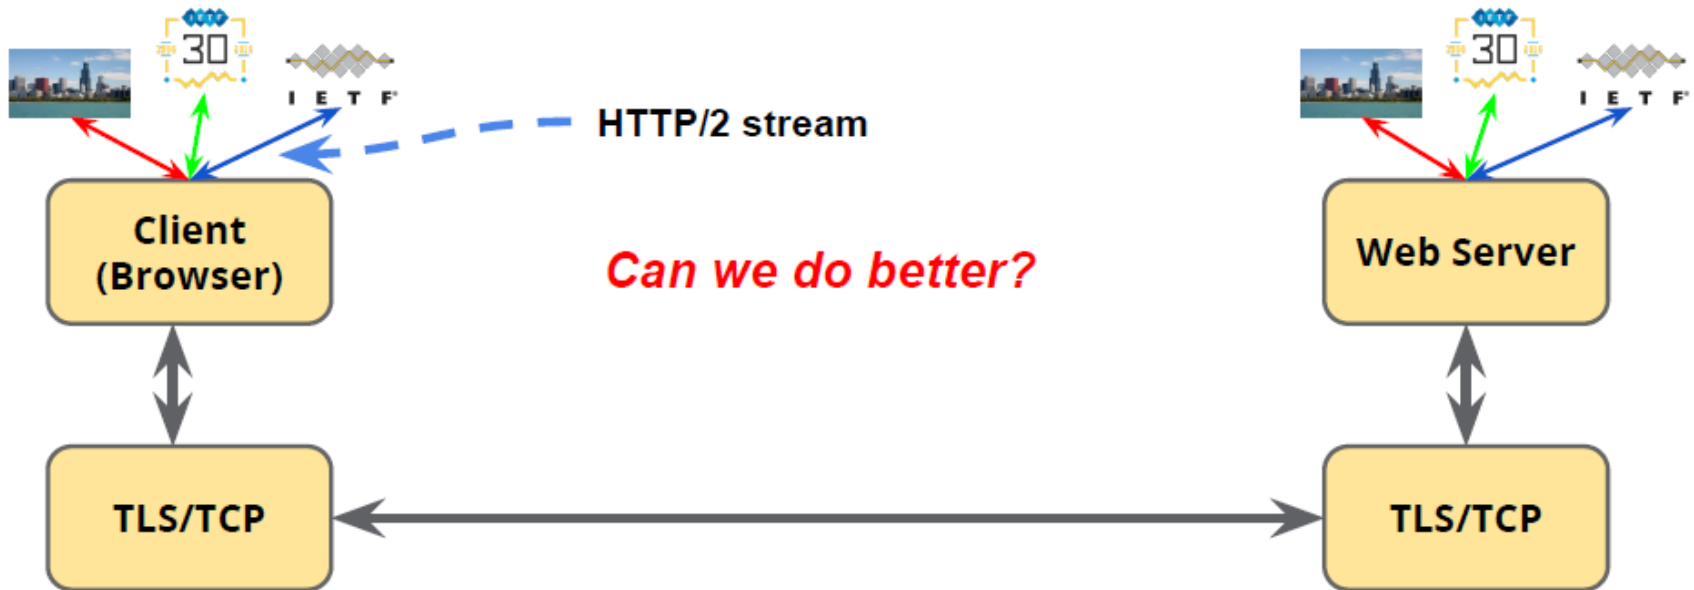
Dealing with Head of Line (HoL) with HTTP/2



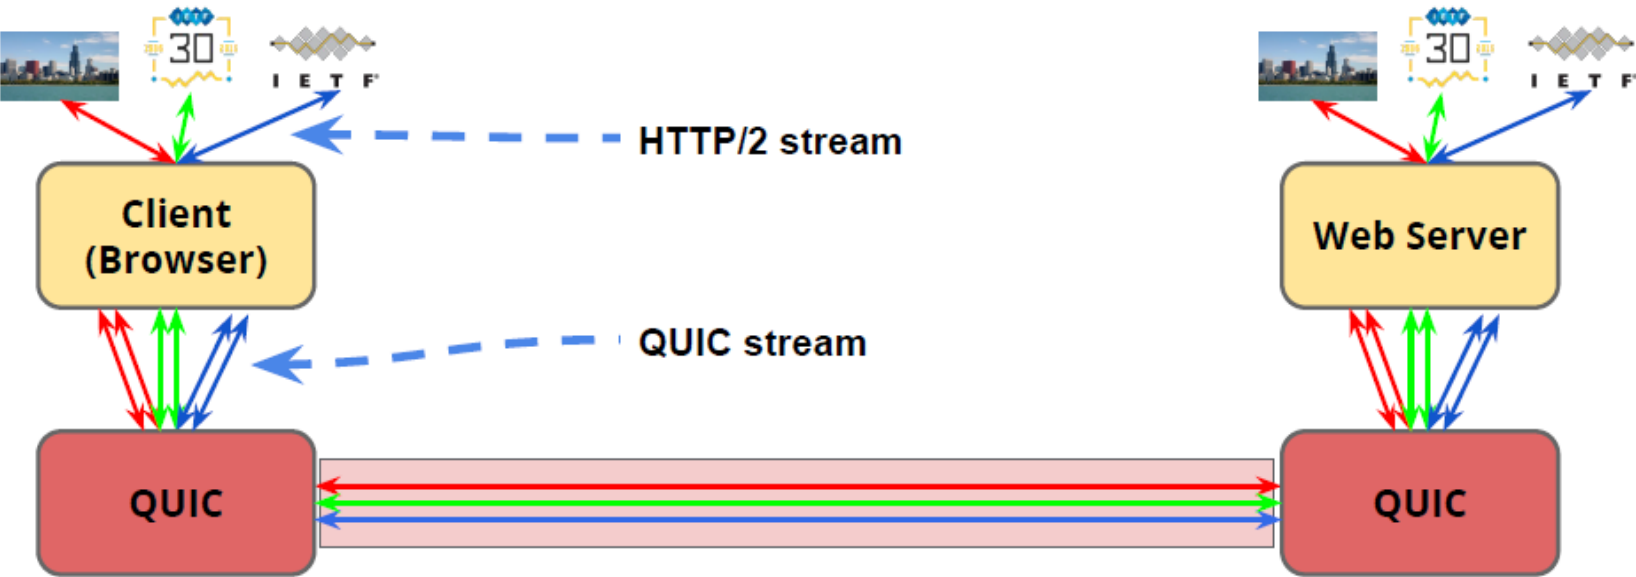
Dealing with Head of Line (HoL) with HTTP/1.1



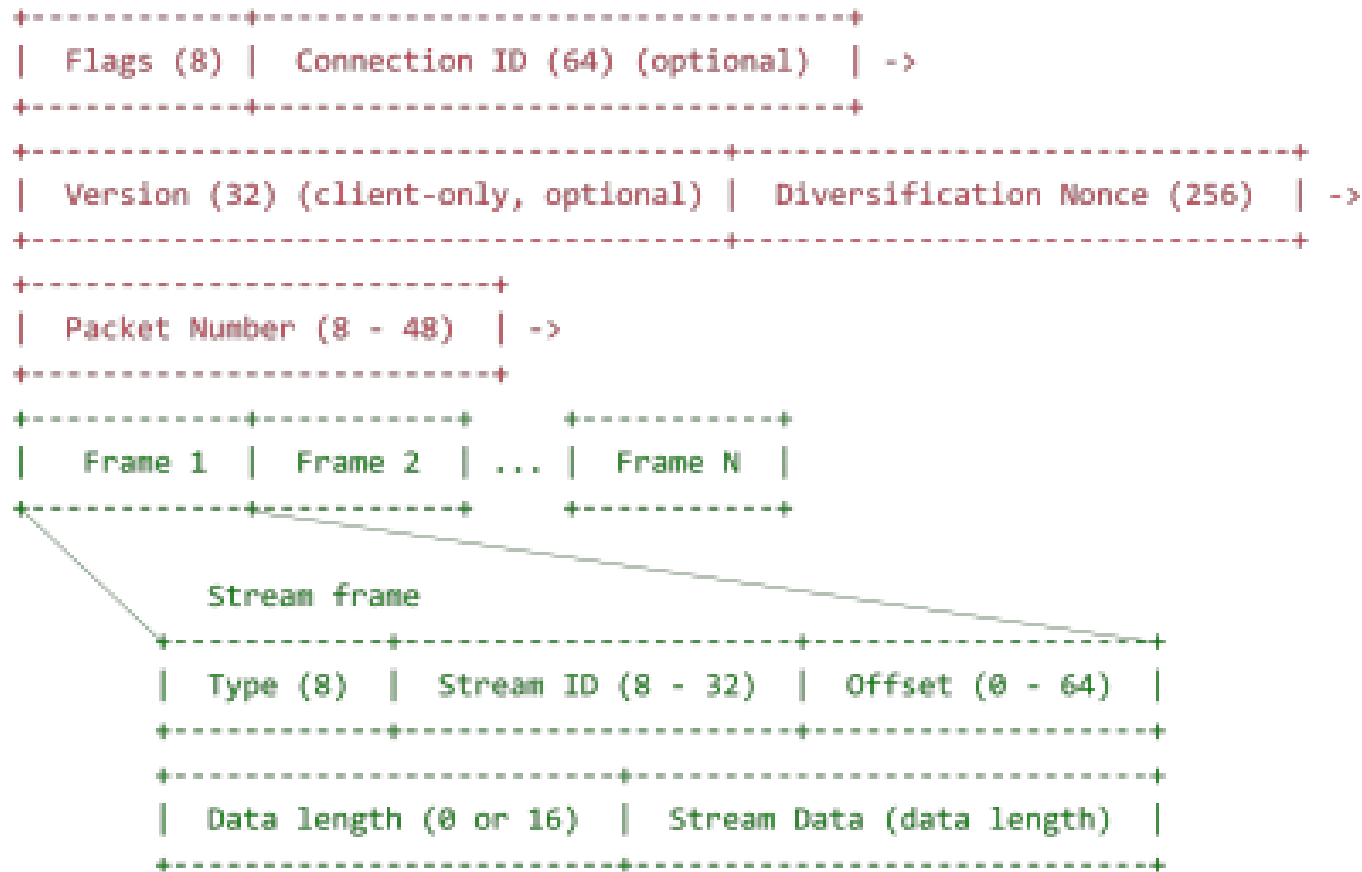
Dealing with Head of Line (HoL) with HTTP/2



HTTP over QUIC



Structure of QUIC Packet



QUIC Support

- **Client**
 - Chrome enable by default
 - Wireshark support
- **Library**
 - libquic / goquic
 - proto-quic
 - First release 4/1
 - Supported by Google

Further References

- IETF draft
 - <http://tools.ietf.org/html/draft-tsvwg-quick-protocol-01>
- www.chromium.org/quick

TCP Congestion Control

- TCP uses **loss-based congestion control strategy**
 - Poor performance with high BW links and large buffer sizes
 - Large buffers leads to long RTTs and delayed congestion notification


BBR Congestion Control

- BBR provides a “**queueless**” congestion control
- A flow should ideally have data in-flight equal to Bandwidth Delay Product (BDP)
 - $BDP = RT_{prop} \times Bt/Bw$
 - At this point, a connection completely saturates the bottleneck link while maintaining an empty buffer
- **What is congestion...?**
 - In-flight data more than BDP (for a long duration) considered as congestion
- **Reference**
 - Congestion Based Congestion Control- BBR [N Cardwell 2016]

BBR



- Essential Path characteristics for congestion control
 - RT_{prop} and Bt/Bw
- **RT_{prop}** and **Bt/Bw** cannot be measured simultaneously. Why?
 - Measuring RT_{prop} requires operating to the left of BDP while measuring Bt/Bw requires operation to the right.
- Both parameters are independent...???



Characterizing the Bottleneck

- A connection runs with the highest throughput and lowest delay when
 - a) Bottleneck packet arrival rate equals Bt/Bw (**rate balance**)
 - b) Total data in-flight equals to $BDP = RT_{prop} \times Bt/Bw$ (**full pipe**)
- It gives 100% utilization with no overflow
- Both conditions should meet simultaneously to ensure “**no queue**”

BtlBw Estimation [.1]

- Average delivery rate between send and ack
 - $deliveryRate = \Delta delivered / \Delta t$
- This rate must be \leq the bottleneck rate
 - Arrival amount is known exactly so all the uncertainty associated with Δt ($\Delta t \geq$ true arrival interval)
 - Therefore, the ratio must be \leq the true delivery rate and upper-bounded by the bottleneck capacity
- BtlBw is given as Windowed-max of delivery rate
$$\widehat{BtlBw} = \max(deliveryRate_t) \quad \forall t \in [T - W_B, T]$$
 - where the time window W_B is typically six to ten RTTs.

BtlBw Estimation [..2]

- TCP must record the departure time of each packet to compute RTT
- BBR augments that record with the total data delivered so each ack arrival yields both an RTT and a delivery rate measurement
- Objective of BBR
 - Matching the packet flow to the delivery path

BBR Algorithm: When ack is received



- Two parts
 - When ack is received
 - When data is sent

```
function onAck(packet)
    rtt = now - packet.sendtime
    update_min_filter(RTpropFilter, rtt)
    delivered += packet.size
    delivered_time = now
    deliveryRate = (delivered - packet.delivered)
                  /(now - packet.delivered_time)
    if (deliveryRate > BtlBwFilter.currentMax
        || ! packet.app_limited)
        update_max_filter(BtlBwFilter,
                          deliveryRate)
    if (app_limited_until > 0)
        app_limited_until -= packet.size
```


BBR Algorithm: When data is sent

```
function send(packet)
    bdp = BtlBwFilter.currentMax
        * RTpropFilter.currentMin
    if (inflight >= cwnd_gain * bdp)
        // wait for ack or timeout
        return
    if (now >= nextSendTime)
        packet = nextPacketToSend()
        if (! packet)
            app_limited_until = inflight
            return
        packet.app_limited =
            (app_limited_until > 0)
        packet.sendtime = now
        packet.delivered = delivered
        packet.delivered_time = delivered_time
        ship(packet)
        nextSendTime = now + packet.size /
            (pacing_gain *
             BtlBwFilter.currentMax)
    timerCallbackAt(send, nextSendTime)
```

Next Topic...

- Congestion Control at Routers-Queuing Algorithms
 - Fair Queuing (FQ)
 - Nagle's FQ Algorithm
 - Max-Min Fairness
 - Weighted Fair Queuing (WFQ)
 - Other Queuing Algorithms (FIFO, CSFQ, RED)
- Reading
 - Random Early Detection Gateways for Congestion Avoidance by Sally Floyd 1993

Thank You!