



BITS Pilani
Pilani Campus

Advance Computer Networks (CS G525)

Virendra S Shekhawat
Department of Computer Science and Information Systems



BITS Pilani
Pilani Campus



First Semester 2018-2019

Slide_Deck_M3_3

Agenda



- Congestion Control at Routers-Queuing Algorithms
 - Fair Queuing (FQ)
 - Nagle's FQ Algorithm
 - Max-Min Fairness
 - Weighted Fair Queuing (WFQ)
 - Other Queuing Algorithms (FIFO, CSFQ, RED)
- Reading
 - Random Early Detection Gateways for Congestion Avoidance by Sally Floyd 1993

Congestion Control: Queuing Algorithms

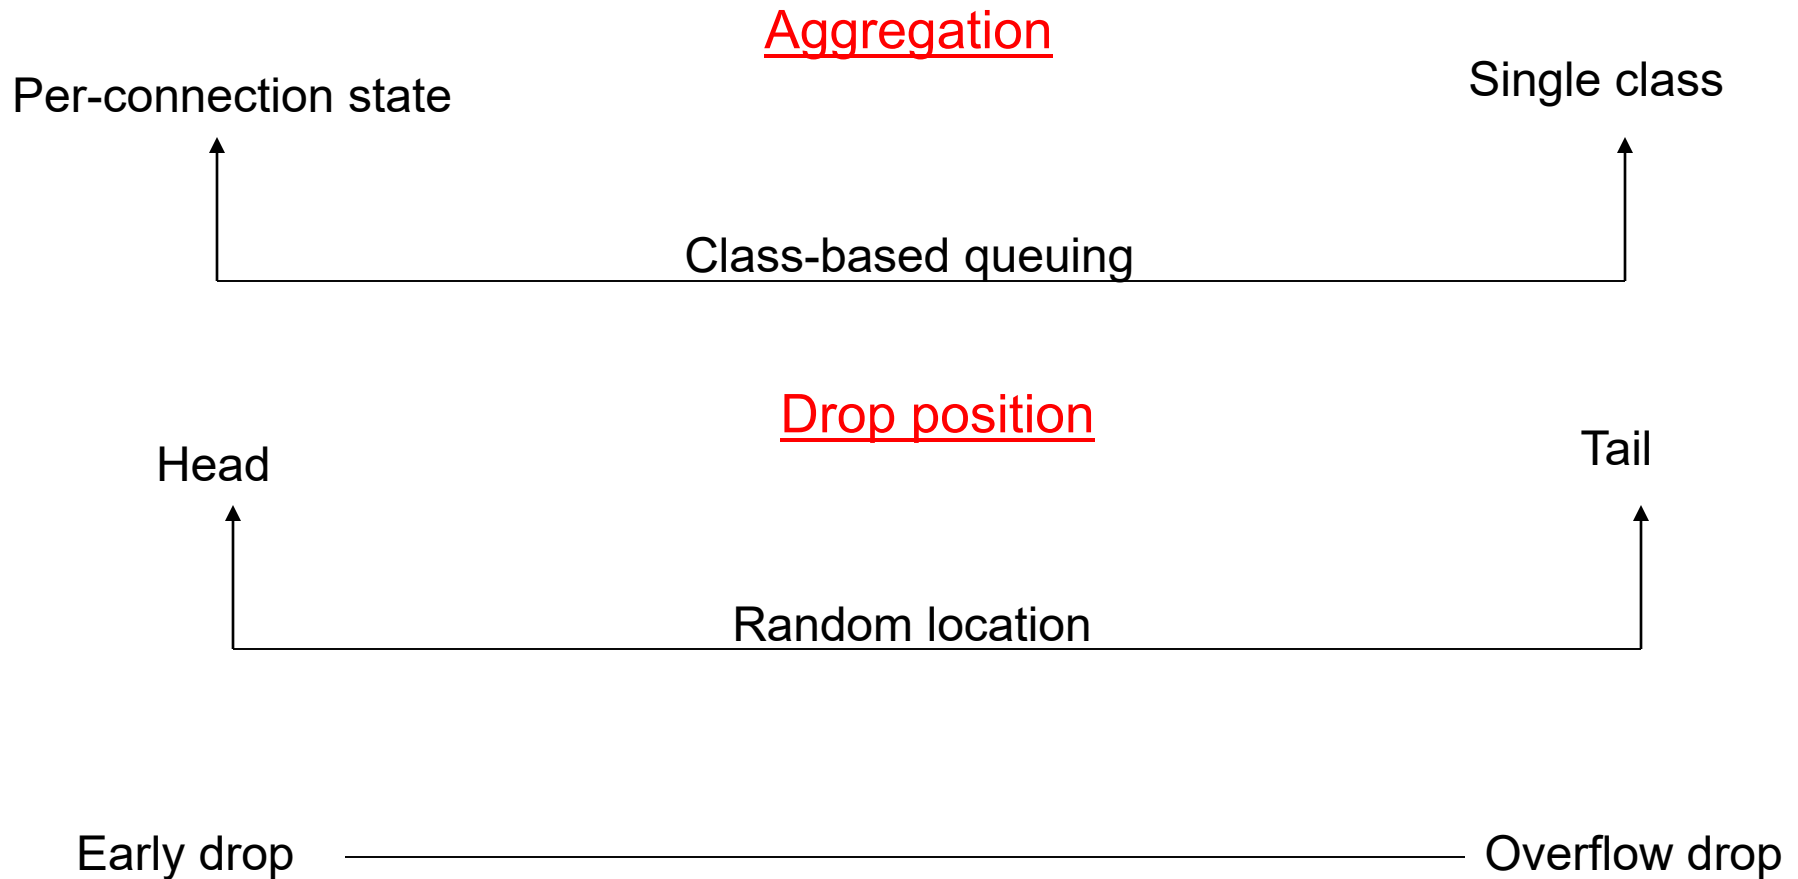


- Congestion can be controlled at gateways through routing and queuing algorithms
- What does queuing algorithms...?
 - Controls the order in which packets are sent
 - Usages of the gateways buffer space
 - Important:
 - Do not change the total traffic on gateways outgoing links
 - Affects the collective behavior of flow control algorithms

Queuing Algorithms Functions

- Which packet get transmitted? (**Bandwidth**)
- When do these packets get transmitted? (**Promptness**)
- Which and when packets get discarded by the gateway? (**Buffer Space**)
- Queuing also affects the latency....!

Packet Drop Dimensions



Fairness among Flows

- At what granularity?
 - s-d pair, source base, receiver base, process base?
- What if users have different RTTs/links/etc.
 - Should it share a link fairly or be TCP fair?
- Maximize fairness index?
 - Fairness = $(\sum x_i)^2 / n(\sum x_i^2)$ $0 < \text{fairness} < 1$

Fairness Goals For Routers

- Allocate resources fairly
 - Isolate ill-behaved users
- Router does not send explicit feedback to source
 - Still needs e2e congestion control
- Should achieve statistical muxing
 - One flow can fill entire pipe if no contenders
 - Work conserving → scheduler never idles link if it has a packet

Nagle's FQ Algorithm

- Gateway maintains separate queues for packets from each individual source
 - Queues are serviced in Round Robin manner
 - Prevents source from sending packet too quickly
 - It increases its own queue length only!
- Drawback
 - Packet size was not considered
 - Source using large packets get more bandwidth

Max-min Fairness

- Allocate user with “small” demand what it wants, evenly divide unused resources to “big” users
- Formally:
 - a) Resources allocated in terms of increasing demand
 - b) No source gets resource share larger than its demand
 - c) Sources with unsatisfied demands get equal share of resource

Max-min Fairness Algorithm

- Assume sources $1....n$, with resource demands $X_1....X_n$ in ascending order
- Assume channel capacity is C
 - Give C/n to X_1 ; if this is more than X_1 wants, divide excess $(C/n - X_1)$ to other sources: each gets $C/n + (C/n - X_1)/(n-1)$
 - If this is larger than what X_2 wants, repeat the process

How to Implement max-min Fairness...?



- Generalized processor sharing
 - Fluid fairness
 - Bitwise round robin among all queues
 - Practical feasibility...?
- Why not simple round robin?
 - Variable packet length → can get more service by sending bigger packets
 - Unfair instantaneous service rate
 - Packets arrive just before/after packet departs?

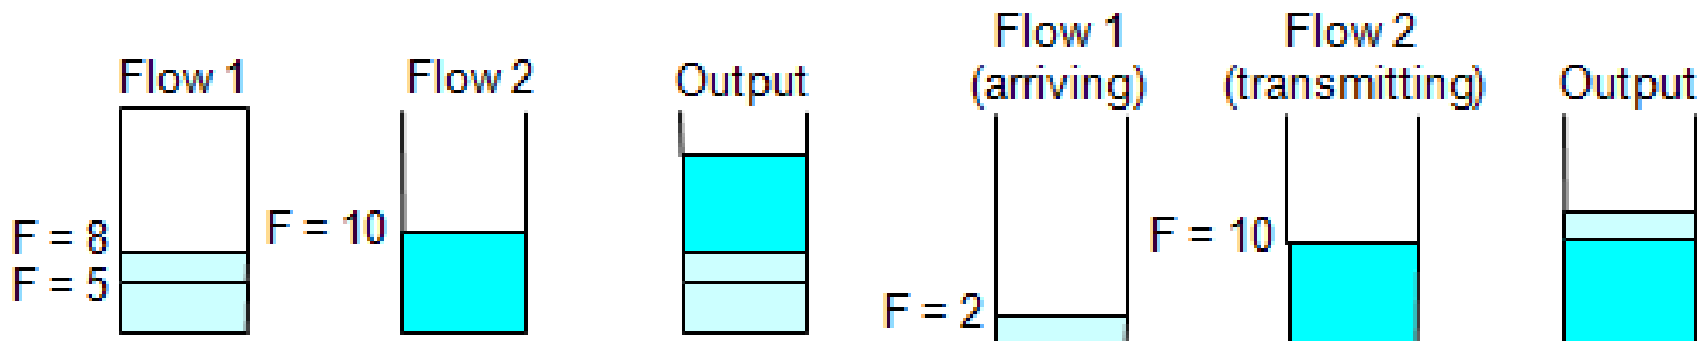
FQ Algorithm



- Suppose clock ticks each time a bit is transmitted
- Let P_i denote the length of packet i
- Let S_i denote the time when start to transmit packet i
- Let F_i denote the time when finish transmitting packet i
- $F_i = S_i + P_i$
- When does router start transmitting packet i ?
 - If packet i arrives before router finished packet $i - 1$ from this flow, then immediately after last bit of $i - 1$ (F_{i-1})
 - If no current packets for this flow, then start transmitting when arrives (call this A_i)
- Thus: $F_i = \text{MAX} (F_{i-1}, A_i) + P_i$

FQ Algorithm (cont...)

- For multiple flows
 - Calculate F_i for each packet that arrives on each flow
 - Treat all F_i 's as time stamps
 - Next packet to transmit is one with lowest timestamp
- Not perfect: can't preempt current packet
- Example



Fair Queuing Tradeoffs



- FQ can control congestion by monitoring flows
 - Non-adaptive flows can still be a problem – why?
- Complex state
 - Must keep queue per flow
 - Hard in routers with many flows (e.g., backbone routers)
 - Flow aggregation is a possibility (e.g. do fairness per domain)
- Complex computation
 - Classification into flows may be hard
 - Must keep queues sorted by finish times
 - Finish times change whenever the flow count changes

Core-Stateless Fair Queuing

- Key problem with FQ is core routers
 - Must maintain state for 1000's of flows
 - Must update state at Gbps line speeds
- CSFQ (Core-Stateless FQ) objectives
 - Edge routers should do complex tasks since they have fewer flows
 - Maintains per flow state
 - Core routers can do simple tasks
 - Core routers can only decide on dropping packets based on the level of congestion

Core-Stateless Fair Queuing

- **Edge Router**
 - Maintains Flow per state.
 - Label the packet with the rate estimate
- **Core Router**
 - Uses FIFO Queuing
 - Probabilistic dropping depending on label, fair rate of router depending on aggregate traffic



Edge Router Behavior

- Monitor each flow i to measure its arrival rate (r_i)

- EWMA of rate
$$r_i^{new} = (1 - e^{-T_i^k/K}) \frac{l_i^k}{T_i^k} + e^{-T_i^k/K} r_i^{old}$$

Here
$$T_i^k = t_i^k - t_i^{k-1}$$

- Non-constant EWMA constant
 - $e^{-T/K}$ where T = current inter-arrival, K = constant
 - Helps to adapt different packet sizes and arrival patterns
- Rate is attached to each packet

Core Router Behavior

- Keep track of fair share rate α
 - $F(\alpha) = \sum_i \min(r_i, \alpha) \rightarrow$ what does this look like?
 - Periodically update α
 - Keep track of current arrival rate
 - Only update α if entire period was congested or uncongested
- Drop probability for packet = $\max(1 - \alpha/r, 0)$

FIFO + DropTail

- Widely used in the Internet
- FIFO (first-in-first-out)
 - Implies single class of traffic
- Drop-tail
 - Arriving packets get dropped when queue is full regardless of flow or importance
- Important distinction:
 - FIFO: scheduling discipline
 - Drop-tail: drop policy

FIFO + Drop-tail Problems

- Leaves responsibility of congestion control to the edges (e.g., TCP)
- Does not separate between different flows
- No policing: send more packets → get more service
- **Synchronization Problem**
 - When queue overflow, packet dropped by gateway for almost all flows hence hosts react to same events

Internet Queuing Problems

- Full queues
 - Routers are forced to have large queues to maintain high utilizations
 - TCP detects congestion from loss
 - Forces network to have long standing queues in steady-state which increases average delay in the network
- Lock-out problem
 - Drop-tail routers treat bursty traffic poorly
 - Traffic gets synchronized easily → allows a few flows to monopolize the queue space (prevent new flows to enter the queue)

Lock-out Problem: Solution

- **Random drop**
 - Packet arriving when queue is full causes some random packet to be dropped
- **Drop front**
 - On full queue, drop packet at head of queue
- **Random drop and drop front solve the lock-out problem but not the full-queues problem**

Full Queues Problem: Solution

- Drop packets before queue becomes full (early drop)
- ***Intuition:*** notify senders of incipient congestion
- Example: Early Random Drop (ERD):
 - If queue length $>$ drop level, drop each new packet with fixed probability p
 - Will this control misbehaving users...?

Active Queue Management (AQM)



- Design active router queue management to aid congestion control
- Why?
 - Routers can distinguish between propagation and persistent queuing delays
 - Routers can decide on transient congestion, based on workload

Active Queue Designs

- Solutions that requires to modify both router and hosts
 - DECbit: congestion bit in packet header
 - Calculates the **avg queue length** for last cycle (busy+idle) plus the current busy period
 - When avg queue length exceeds → set the congestion bit in the header of the arriving packets
 - If less than 50% of last window's worth had bit set
 - Increase `congestionWindow` by 1 packet
 - If 50% or more of last window's worth had bit set
 - Decrease `congestionWindow` by 0.875 times
- Solutions that requires to modify only router and host uses TCP
 - Fair queuing
 - Per-connection buffer allocation
 - RED (Random Early Detection)
 - Drop packet or set bit in packet header as soon as congestion is starting

Design Objectives: Queuing Algorithms



- Keep throughput high and delay low
- Accommodate bursts to avoid bias against bursty traffic
- Queue size should reflect ability to accept bursts rather than steady-state queuing
- Maintain an upper bound on the average queue length even in the absence of cooperation from TP protocols
- Avoidance of global synchronization

RED Algorithm

- Maintain running average of queue length
- If $\text{avgq} < \text{min}_{\text{th}}$ do nothing
 - Low queuing, send packets through
- If $\text{avgq} > \text{max}_{\text{th}}$, drop packet
 - Protection from misbehaving sources
- Else mark packet in a manner proportional to queue length
 - Notify sources of incipient congestion
- Maintain running average of queue length
 - Byte mode vs. packet mode – why?

RED Algorithms

- Two separate algorithms
 - One for computing the average queue size
 - To determine the degree of burstiness
 - Second for calculating packet marking probability
 - Gives current level of congestion
 - More packet are marked indicate higher congestion

Queue Estimation

- Standard EWMA: $avgq = (1-w_q) avgq + w_q qlen$
 - Special fix for idle periods (empty queue)– why?
 - Assume m packets of small size processed during idle period
- Upper bound on w_q depends on min_{th}
 - Want to ignore transient congestion
 - Set w_q such that certain burst size does not exceed min_{th}
- Lower bound on w_q is to detect congestion relatively quickly
- Typical $w_q = 0.002$

Thresholds

- \min_{th} determined by the utilization requirement
 - Tradeoff between queuing delay and utilization
- Relationship between \max_{th} and \min_{th}
 - Want to ensure that feedback has enough time to make difference in load
 - Depends on average queue increase in one RTT
 - Paper suggests ratio of two

Packet Marking

- Marking probability is based on queue length
 - $P_b = \max_p(\text{avgq} - \min_{th}) / (\max_{th} - \min_{th})$
- Just marking based on P_b can lead to clustered marking
 - Could result in synchronization
 - Better to bias P_b by history of unmarked packets
 - $P_a = P_b / (1 - \text{count} * P_b)$
 - This ensures gateway doesn't wait too long to mark a packet

RED Algorithm

Initialization:

$avg \leftarrow 0$

$count \leftarrow -1$

for each packet arrival

calculate new avg. queue size avg :

if the queue is nonempty

$avg \leftarrow (1 - w_q)avg + w_q q$

else

$m \leftarrow f(time - q_time)$

$avg \leftarrow (1 - w_q)^m avg$

if $min_{th} \leq avg < max_{th}$

increment $count$

calculate probability p_a :

$p_b \leftarrow$

$max_p(avg - min_{th}) / (max_{th} - min_{th})$

$p_a \leftarrow p_b / (1 - count \cdot p_b)$

with probability p_a :

mark the arriving packet

$count \leftarrow 0$

else if $max_{th} \leq avg$

mark the arriving packet

$count \leftarrow 0$

else $count \leftarrow -1$

when queue becomes empty

$q_time \leftarrow time$

RED - Summary

- Detect incipient congestion, allow bursts
- Keeps power (throughput/delay) high
 - Keep average queue size low
 - Assume hosts respond to lost packets
- Avoids window synchronization
 - Randomly mark packets
- Avoids bias against bursty traffic
- Some protection against ill-behaved users

Extending RED for Flow Isolation



- Problem: what to do with non-cooperative flows?
- Fair queuing achieves isolation using per-flow state – expensive at backbone routers
 - How can we isolate unresponsive flows without per-flow state?
- RED penalty box
 - Monitor history for packet drops, identify flows that use disproportionate bandwidth
 - Isolate and punish those flows

Thank You!